

```

import pandas as pd
import numpy as np
from collections import defaultdict, Counter
from datasets import load_dataset, Dataset as HFDataset
from sklearn.metrics import accuracy_score
import warnings
import string
import os
import shutil
import torch
import re
import random
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import sentencepiece
import evaluate
from sacrebleu import BLEU
from transformers import AutoTokenizer, AutoModelForSeq2SeqLM, Seq2SeqTrainingArguments, Seq2SeqTrainer, pipeline, EvalPrediction, Traine
import sys
import tensorflow as tf

```

▼ Data Preprocessing

▼ Load The Data

```

dataset = load_dataset(" ")
print(dataset.keys())

```

```
dict_keys(['train', 'test'])
```

```

print(dataset['train'].column_names)
print(dataset['train'][0])

```

```

['index', 'Text', 'label']
{'index': 2835, 'Text': 'مائل ذاته ليس غير.. لو قدر لتبعة الماء أن تحس، إذن لأحسب ذلك الشعور، العطاء المحض الذي يخلق من جديد كلما شرب غير من مائها'

```

```

df = dataset['train'].to_pandas()
df.shape

```

```
(2380, 3)
```

▼ Clean and Normalize Words

```

def normalize_arabic(text):
    text = re.sub("[\u0600]", "", text)
    text = re.sub("ى", "ي", text)
    text = re.sub("ؤ", "أ", text)
    text = re.sub("ئ", "أ", text)
    text = re.sub("ة", "ا", text)
    text = re.sub("ك", "ك", text)
    text = re.sub(" و ", " ", text)
    return text

```

```

# Define Arabic diacritics
arabic_diacritics = re.compile(r"""
' | # Shadda
' | # Fatha
' | # Tanwin Fath
' | # Damma
' | # Tanwin Damm
' | # Kasra
' | # Tanwin Kasr
' | # Sukun
- # Tatweel
""", re.VERBOSE)

```

```

def remove_diacritics(text):
    text = re.sub(arabic_diacritics, '', text)
    return text

```

```
def remove_non_arabic_and_numbers(text):
```

```

# Keep only Arabic characters and spaces
arabic_pattern = re.compile(r'^\u0600-\u06FF\u0750-\u077F\u08A0-\u08FF\uFB50-\uFDFD\uFE70-\uFEFF\s$')
return arabic_pattern.sub('', text)

# Define punctuation characters
arabic_punctuations = ''`÷×−“”‘’!|+|~{}',.!":/~][%^&*()_<>¡''
english_punctuations = string.punctuation
punctuations = arabic_punctuations + english_punctuations

def remove_punctuations(text):
    translator = str.maketrans('', '', punctuations)
    return text.translate(translator)

def normalize_arabic_text(text):
    text = normalize_arabic(text)
    text = remove_diacritics(text)
    text = remove_non_arabic_and_numbers(text)
    text = remove_punctuations(text)
    return text

# Extract all text from the 'Text' column
text = ' '.join(df["Text"])
text = normalize_arabic_text(text)
words = text.split()
vocab = list(set(words))
print(len(vocab))

```

 72804

▼ Generate Synthetic Training Data

▼ Insertion

```

def insert(word, max_typos=2):
    arabic_letters = 'اَبْثَجْخْذْزْسْشْصْضْطْظْعْفْقْكْلْمْنْهْوْي'
    possible_words = set()
    for _ in range(max_typos):
        pos = random.randint(0, len(word))
        letter = random.choice(arabic_letters)
        new_word = word[:pos] + letter + word[pos:]
        possible_words.add(new_word)
    return possible_words

```

▼ Deletion

```

def delete(word):
    if len(word) < 2:
        return set()
    pos = random.randint(0, len(word)-1)
    return {word[:pos] + word[pos+1:]}

```

▼ Replacing

```

def replace(word):
    arabic_letters = 'اَبْثَجْخْذْزْسْشْصْضْطْظْعْفْقْكْلْمْنْهْوْي'
    if len(word) < 1:
        return set()
    pos = random.randint(0, len(word)-1)
    letter = random.choice(arabic_letters)
    return {word[:pos] + letter + word[pos+1:]}

```

▼ Switching

```

def switch(word):
    if len(word) < 2:
        return set()
    pos = random.randint(0, len(word)-2)
    new_word = word[:pos] + word[pos+1] + word[pos] + word[pos+2:]
    return {new_word}

```

```
def add_arabic_typo(word, max_total_typos=3):
    if len(word) < 2:
        return [word]

    common_typos = {
        'ا': ['ا', 'ا', 'ا', 'ا'],
        'إ': ['إ', 'إ', 'ء'],
        'ة': ['ة', 'ت'],
        'ه': ['ه', 'ح'],
        'ي': ['ى', 'ئ'],
        'و': ['و', 'ء', ''],
        'ن': ['ت', 'ب'],
        'ت': ['ث', 'ن'],
        'س': ['ش', 'ص'],
    }

    typo_ops = [insert, delete, replace, switch]
    typo_variants = set()

    for _ in range(2):
        pos = random.randint(0, len(word)-1)
        char = word[pos]
        if char in common_typos:
            replacement = random.choice(common_typos[char])
            typo = word[:pos] + replacement + word[pos+1:]
            typo_variants.add(typo)

    while len(typo_variants) < max_total_typos:
        op = random.choice(typo_ops)
        typo_variants.update(op(word))

    return list(typo_variants)
```

▼ Create Training Pairs

```
synthetic_data = []

for word in vocab[:10000]:
    typo_variants = add_arabic_typo(word)
    for incorrect in typo_variants:
        if incorrect != word:
            synthetic_data.append({'inputs': incorrect.strip(), 'targets': word.strip()})

df = pd.DataFrame(synthetic_data)

df.to_csv("autocorrect_dataset.csv", index=False)

with open("autocorrect_dataset.csv", encoding="utf-8") as f:
    for i in range(10):
        print(f.readline())
```

→ inputs, targets

الاسرار والاسرار

الاسرار والاسرار

الاسرار والاسرار

الاسرار والاسرار

بومركيه وبومركيه

مركيه وبومركيه

زومركيه وبومركيه

باتوا وباتوا

باتوا وباتوا

df.shape

→ (31851, 2)

✓ Data Cleaning

```
def clean_dataset(df):

    df = df.dropna(subset=["inputs", "targets"])

    df = df[df["inputs"].apply(lambda x: isinstance(x, str))]
    df = df[df["targets"].apply(lambda x: isinstance(x, str))]

    df = df.reset_index(drop=True)

    return df
```

✓ Data Splitting

```
# Load and clean dataset
df = pd.read_csv("autocorrect_dataset.csv")
df = clean_dataset(df)

# Train/validation split
train_df, val_df = train_test_split(df, test_size=0.1, random_state=42)
train_dataset = HFDataset.from_pandas(train_df)
val_dataset = HFDataset.from_pandas(val_df)
```

```
df.shape
```

```
(31851, 2)
```

✓ Check Duplicated Values

```
duplicates = train_df[train_df['inputs'] == train_df['targets']]
print(f"Found {len(duplicates)} identical input-output pairs")
```

```
Found 0 identical input-output pairs
```

✓ Model Fine-Tuning

✓ Load Tokenizer & Model

```
tokenizer = AutoTokenizer.from_pretrained("CAMEL-Lab/arat5-coda", legacy=False)
model = AutoModelForSeq2SeqLM.from_pretrained("CAMEL-Lab/arat5-coda")

def tokenize(batch):
    input_encodings = tokenizer(batch['inputs'], padding="max_length", truncation=True, max_length=128)
    target_encodings = tokenizer(batch['targets'], padding="max_length", truncation=True, max_length=128)

    labels = target_encodings["input_ids"]
    labels = [
        [(1 if l != tokenizer.pad_token_id else -100) for l in label]
        for label in labels
    ]

    return {
        "input_ids": input_encodings["input_ids"],
        "attention_mask": input_encodings["attention_mask"],
        "labels": labels
    }
```

```
# Apply tokenization
train_dataset = train_dataset.map(tokenize, batched=True, remove_columns=["inputs", "targets"])
val_dataset = val_dataset.map(tokenize, batched=True, remove_columns=["inputs", "targets"])
```

```
Map: 100%|██████████| 28665/28665 [00:16<00:00, 1705.89 examples/s]
Map: 100%|██████████| 3186/3186 [00:01<00:00, 1660.88 examples/s]
```

✓ Model Training Setup

```

class LossTrackerCallback(TrainerCallback):
    def __init__(self):
        self.eval_losses = []
        self.train_losses = []

    def on_evaluate(self, args, state, control, metrics, **kwargs):
        if "eval_loss" in metrics:
            self.eval_losses.append((state.global_step, metrics["eval_loss"]))

    def on_log(self, args, state, control, logs=None, **kwargs):
        if logs and "loss" in logs:
            self.train_losses.append((state.global_step, logs["loss"]))

training_args = Seq2SeqTrainingArguments(

    output_dir="./results",
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    warmup_steps=200,
    warmup_ratio=0.1,
    weight_decay=0.01,
    gradient_accumulation_steps=1,
    learning_rate=2e-5,
    fp16=True,
    logging_dir="./logs",
    logging_steps=500,
    predict_with_generate=True,
    generation_max_length=16,
    gradient_checkpointing=True,
    max_grad_norm=1.0,
    lr_scheduler_type="cosine",
    optim="adafactor",
    report_to="tensorboard"
)

loss_tracker = LossTrackerCallback()

# Trainer
trainer = Seq2SeqTrainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    tokenizer=tokenizer,
    compute_metrics=compute_metrics,
    callbacks=[loss_tracker]
)

```

✓ Model Training

✓ Train The Model

```

num_epochs = 2

for epoch in range(num_epochs):
    print(f"\nEpoch {epoch + 1}/{num_epochs}")
    trainer.train()
    eval_results = trainer.evaluate()
    print(f"Validation Loss: {eval_results['eval_loss']:.4f}")

```



Epoch 1/2

`use_cache=True` is incompatible with gradient checkpointing. Setting `use_cache=False`...

[5376/5376 13:51:36, Epoch 3/3]

Step	Training Loss
------	---------------

500	3.896200
-----	----------

1000	2.641700
------	----------

1500	2.358500
------	----------

2000	2.220500
------	----------

2500	2.088500
------	----------

3000	2.054900
------	----------

3500	2.009200
------	----------

4000	1.913900
------	----------

4500	1.923300
------	----------

5000	1.894700
------	----------

Passing a tuple of `past_key_values` is deprecated and will be removed in Transformers v4.48.0. You should pass an instance of `Enc

[200/200 10:52:16]

Metric computation error: name 'bleu_scorer' is not defined

Validation Loss: 1.5322

Epoch 2/2

[5376/5376 10:38:22, Epoch 3/3]

Step	Training Loss
------	---------------

500	1.913600
-----	----------

1000	1.863400
------	----------

1500	1.792000
------	----------

2000	1.737900
------	----------

2500	1.659600
------	----------

3000	1.646200
------	----------

3500	1.649300
------	----------

4000	1.556200
------	----------

4500	1.583800
------	----------

5000	1.555700
------	----------

Metric computation error: name 'bleu_scorer' is not defined

Validation Loss: 1.3276

✓ Evaluation

```
from datasets import load_dataset
from transformers import AutoTokenizer, AutoModelForSequenceClassification
import torch
import evaluate
```

```
label_names = list(set(dataset["train"]["label"]))
label2id = {label: i for i, label in enumerate(label_names)}
id2label = {i: label for label, i in label2id.items()}
```

```
# Load tokenizer and model
model_path = "D:\\NLP_model\\my_model"
tokenizer = AutoTokenizer.from_pretrained(model_path)
model = AutoModelForSequenceClassification.from_pretrained(model_path)
```

```
# Load metric
metric = evaluate.load("accuracy")
```

```
# Evaluation loop
total_samples = 100
predictions = []
references = []
```

```
for example in dataset["test"].select(range(total_samples)):
    text = example["Text"]
    true_label = label2id[example["label"]] # convert to integer
```

```

inputs = tokenizer(text, return_tensors="pt", truncation=True, padding=True)
with torch.no_grad():
    outputs = model(**inputs)
    pred_id = torch.argmax(outputs.logits, dim=1).item()

predictions.append(pred_id)
references.append(true_label)

# Compute accuracy
results = metric.compute(predictions=predictions, references=references)
print(f"\n Accuracy on {total_samples} test samples: {results['accuracy']:.2f}")

```

⚠ Some weights of T5ForSequenceClassification were not initialized from the model checkpoint at D:\NLP_model\my_model and are newly in
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Accuracy on 100 test samples: 0.05

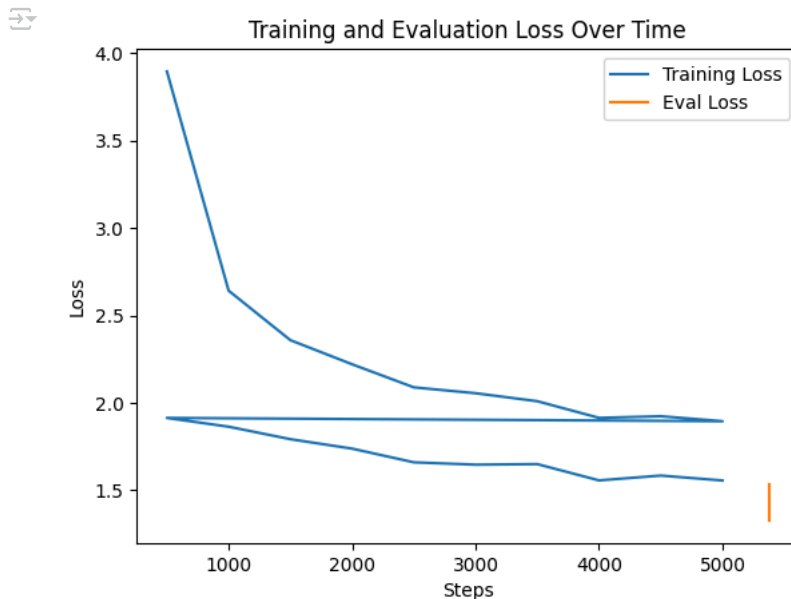
```

# Plot training loss
steps, train_losses = zip(*loss_tracker.train_losses)
plt.plot(steps, train_losses, label="Training Loss")

# Plot evaluation loss
if loss_tracker.eval_losses:
    eval_steps, eval_losses = zip(*loss_tracker.eval_losses)
    plt.plot(eval_steps, eval_losses, label="Eval Loss")

plt.xlabel("Steps")
plt.ylabel("Loss")
plt.legend()
plt.title("Training and Evaluation Loss Over Time")
plt.show()

```



GUI

```

import tkinter as tk
from tkinter import ttk
from transformers import AutoTokenizer, AutoModelForSeq2SeqLM
import torch

# Function to perform correction
def correct_text():
    input_text = input_box.get("1.0", tk.END).strip()
    if not input_text:
        output_label.config(text="يرجى إدخال نص")
        return

    inputs = tokenizer(input_text, return_tensors="pt", padding=True, truncation=True)
    with torch.no_grad():
        outputs = model.generate(**inputs, max_new_tokens=128)

```

```
corrected = tokenizer.decode(outputs[0], skip_special_tokens=True)
output_label.config(text=corrected)

# Create GUI window
root = tk.Tk()
root.title("مصحح اللغة العربية")
root.geometry("600x400")
root.configure(bg="#1e1e2f")

# Title
title = tk.Label(root, text="تصحيح النصوص العربية", font=("Segoe UI", 20), bg="#1e1e2f", fg="ffffff")
title.pack(pady=10)

# Input Text Box
input_box = tk.Text(root, font=("Segoe UI", 14), height=6, wrap=tk.WORD, bg="#2c2c3c", fg="ffffff", insertbackground="white", padx=10,
input_box.pack(padx=20, pady=10, fill=tk.BOTH)

# Correct Button
correct_button = ttk.Button(root, text="تصحیح", command=correct_text)
style = ttk.Style()
style.theme_use("clam")
style.configure("TButton", font=("Segoe UI", 12), foreground="white", background="#3a86ff", padding=10)
style.map("TButton", background=[("active", "#265df2")])
correct_button.pack(pady=10)

# Output Label
output_label = tk.Label(root, text="", font=("Segoe UI", 14), wraplength=560, justify="right", bg="#1e1e2f", fg="#00ffb3")
output_label.pack(padx=20, pady=10)

root.mainloop()
```